

Cache Modelling

Euan Scott-Watson

November 17, 2021

1 Question 1

I created the simulation in Python. Please see below the time/hit-ratio graph for the FIFO policy cache with $m=1024$ and $n=65536$; tested with 3 separate time steps.




Figure 1: Graphs for times of 1,000, 5,000 and 10,000 arbitrary time units

From my simulations I determined that the hit ratio for a FIFO Cache was ~ 0.49 (and ~ 0.53 for LRU). Therefore, the throughput of my cache to the storage device is $1 - \text{Hit Ratio}$, i.e. ~ 0.51 (and ~ 0.47 for LRU). My estimates are unbiased as I have run multiple tests, all converging to the same values across the simulations. Moreover, the events are all independent of each other causing no bias in picking the timings for each element. I have modelled the Poisson processes using the Inverse Transform method: $t_k = \frac{-\ln(U(0,1))}{\lambda_k}$ where $\lambda_k = \frac{1}{k+1}$. I filled a diary with timings for each process and iteratively took the one with the smallest time, fetched it from the cache, then re-calculated its inter-arrival time and put it back into the diary. I continuously did this until the current time, t , had reached the time limit, T . I took regular readings of the number of hits and misses every 25 arbitrary time units to plot the graphs above with. See the code in **Appendix**.

2 Question 2

$$Q = \begin{pmatrix} -(\lambda_1 + \lambda_2) & 0 & \lambda_1 & 0 & \lambda_2 & 0 \\ 0 & -(\lambda_1 + \lambda_2) & \lambda_1 & 0 & \lambda_2 & 0 \\ \lambda_0 & 0 & -(\lambda_0 + \lambda_2) & 0 & 0 & \lambda_2 \\ \lambda_0 & 0 & 0 & -(\lambda_0 + \lambda_2) & 0 & \lambda_2 \\ 0 & \lambda_0 & 0 & \lambda_1 & -(\lambda_0 + \lambda_1) & 0 \\ 0 & \lambda_0 & 0 & \lambda_1 & 0 & -(\lambda_0 + \lambda_1) \end{pmatrix}$$

$S = \{(0, 1), (0, 2), (1, 0), (1, 2), (2, 0), (2, 1)\}$

3 Question 3

The total rate out of state s_i is equal to the negative sum of all the rates of reachable states - not including i , i.e. $s_i = -\sum_{j \neq i}^m \lambda_j \forall Q_{i,j} \neq 0$

4 Question 4

$$\lambda_0 = 1; \lambda_1 = 1/2; \lambda_2 = 1/3$$

$$\mathbf{p}Q_c = 0_c$$

$$Q = \begin{pmatrix} -5/6 & 0 & 1/2 & 0 & 1/3 & 0 \\ 0 & -5/6 & 1/2 & 0 & 1/3 & 0 \\ 1 & 0 & -4/3 & 0 & 0 & 1/3 \\ 1 & 0 & 0 & -4/3 & 0 & 1/3 \\ 0 & 1 & 0 & 1/2 & -3/2 & 0 \\ 0 & 1 & 0 & 1/2 & 0 & -3/2 \end{pmatrix}$$

$$\mathbf{p} = 0_3 Q_3^{-1}$$

$$\mathbf{p} = (0 \ 0 \ 1 \ 0 \ 0 \ 0) \begin{pmatrix} -5/6 & 0 & 1 & 0 & 1/3 & 0 \\ 0 & -5/6 & 1 & 0 & 1/3 & 0 \\ 1 & 0 & 1 & 0 & 0 & 1/3 \\ 1 & 0 & 1 & -4/3 & 0 & 1/3 \\ 0 & 1 & 1 & 1/2 & -3/2 & 0 \\ 0 & 1 & 1 & 1/2 & 0 & -3/2 \end{pmatrix}^{-1}$$

$$\mathbf{p} = (18/55 \ 12/55 \ 0.204 \ 0.068 \ 4/33 \ 2/33)$$

$$\text{Sum of rates} = \frac{11}{6}$$

$$p(0 \text{ or } 1) = \frac{9}{11}$$

$$p(0 \text{ or } 2) = \frac{8}{11}$$


$$p(1 \text{ or } 2) = \frac{5}{11}$$

The hit ratio is calculated by multiplying the probability of a state by the probability of getting one of the elements in that state as the next element, i.e. $p((a, b))p(a \text{ or } b)$: $(\frac{18}{55} \times \frac{9}{11}) + (\frac{12}{55} \times \frac{8}{11}) + (0.204 \times \frac{9}{11}) + (0.068 \times \frac{5}{11}) + (\frac{4}{33} \times \frac{8}{11}) + (\frac{2}{33} \times \frac{5}{11}) = 0.740$ which is similar to the value of 0.721 I got from the simulation.

5 Question 5

$$Q = \begin{pmatrix} -1/3 & 0 & 0 & 1/3 & 0 & 0 \\ 0 & -1/2 & 0 & 0 & 0 & 1/2 \\ 0 & 1/3 & -1/3 & 0 & 0 & 0 \\ 0 & 0 & 0 & -1 & 1 & 0 \\ 1/2 & 0 & 0 & 0 & -1/2 & 0 \\ 0 & 0 & 1 & 0 & 0 & -1 \end{pmatrix}$$

Only taking into consideration those reachable from the initial configuration of (0, 1):



We can only reach the states: (0, 1), (2, 0), (1, 2), thus:

$$Q = \begin{pmatrix} -1/3 & 0 & 1/3 \\ 1/2 & -1/2 & 0 \\ 0 & 1 & -1 \end{pmatrix}$$

$$\mathbf{p} = 0_3 Q_3^{-1}$$

$$\mathbf{p} = (0 \ 0 \ 1) \begin{pmatrix} -1/3 & 0 & 1 \\ 1/2 & -1/2 & 1 \\ 0 & 1 & 1 \end{pmatrix}^{-1}$$

$$\mathbf{p} = (1/2 \ 1/3 \ 1/6)$$

The hit ratio is calculated same as Q4: $(\frac{1}{2} \times \frac{9}{11}) + (\frac{1}{6} \times \frac{5}{11}) + (\frac{1}{3} \times \frac{8}{11}) = \frac{8}{11} = 0.727$, which is similar to the value of 0.743 I got from the simulation.

6 Appendix

6.1 Cache.py Code

```
1 class Cache:
2     def __init__(self, capacity, starting_values=None):
3         self.name = "Cache"
4         self.capacity = capacity
5         self.hits = 0
6         self.misses = 0
7         if starting_values:
8             self.cache = starting_values + [-1 for _ in range(self.capacity - len(starting_values))]
9         else:
10            self.cache = [-1 for _ in range(self.capacity)]
11
12     def evict(self, value):
13         pass
14
15     def hit(self, value):
16         pass
17
18     def get(self, value):
19         if value in self.cache:
20             self.hits += 1
21             self.hit(value)
22             return value
23         else:
24             self.misses += 1
25             self.evict(value)
```

6.2 FIFOCache.py Code

```
1 from Cache import Cache
2
3 class FIFOCache(Cache):
4     def __init__(self, capacity, starting_values=None):
5         super().__init__(capacity, starting_values)
6         self.name = "FIFO"
7
8     def evict(self, value):
9         self.cache = self.cache[1:] + [value]
```

6.3 LRUCache.py Code

```
1 from Cache import Cache
2
3 class LRUCache(Cache):
4     def __init__(self, capacity, starting_values=None):
5         super().__init__(capacity, starting_values)
6         self.name = "LRU"
7
8     def evict(self, value):
9         self.cache = [value] + self.cache[:-1]
10
11     def hit(self, value):
12         index = self.cache.index(value)
13         self.cache = [value] + self.cache[:index] + self.cache[index + 1:]
```

6.4 model.py Code

```
1 import sys
2 import time
3 from Cache import Cache
4 from LRUCache import LRUCache
5 from FIFOCache import FIFOCache
6 import random, math, bisect
7
8 import matplotlib.pyplot as plt
9
10
11 # Functions for sampling
12 def rate(k):
13     return 1 / (k + 1)
14
15
16 def inverse_transform(k):
17     l = rate(k)
18     U = random.uniform(0, 1)
19     return -math.log(U) / l
20
21
22 # Get a start time for all potential ks
23 def sample_start(n):
24     timings = []
25     for k in range(n):
26         t = inverse_transform(k)
27         timings.append((t, k))
28     return sorted(timings, key=lambda x: x[0])
29
30
31 def plot_and_save(time_points, hit_ratio):
32     # Plots the graph and saves it as a png
33     ax = plt.axes()
34     ax.plot(time_points, hit_ratio)
35     plt.xlabel("Time (arbitrary units)")
36     plt.ylabel("Hit-Ratio")
37     plt.title("Hit Ratio Graph for an {0} Cache (m={1}, n={2})".format(cache.name, m, n))
38     plt.savefig("graphs/{0}_{1}_{2}_{3}.png".format(m, n, T, cache.name))
39     plt.show()
40
41
42 def model(n: int, cache: Cache, T: int):
43     start = time.time()
44
45     # Arrays for storing data that will be plotted
46     hit_ratio = [0]
47     time_points = [0]
48     # Variable to keep track of current time in simulation
49     t = 0
50
51     # Variables for how often to take measurements of HR
52     rate_of_data_collection = 1
53     next_data_collected = rate_of_data_collection
54
55     diary = sample_start(n)
56     while t < T:
57         # Get the next entry in the diary
58         event = diary.pop(0)
59         t = event[0]
60         k = event[1]
61         cache.get(k)
62         tP = t + inverse_transform(k)
63         bisect.insort(diary, (tP, k))
64
65         # If enough time has passed, collect more graphing data
66         if t // next_data_collected >= 1:
67             time_points.append(t)
68             hit_ratio.append(cache.hits / (cache.hits + cache.misses))
69             next_data_collected += rate_of_data_collection
70
71         # Print the progress, flushing out the previous print
72         print("%.2f percent complete - %.2f seconds elapsed\t Current HR: %.2f" % (
73             round(t / T * 100, 3),
74             round(time.time() - start, 3),
75             round(cache.hits / (cache.hits + cache.misses), 3)
76         ), end='\r')
77         sys.stdout.flush()
78
```

```
79     print("")
80     plot_and_save(time_points, hit_ratio)
81
82
83 if __name__ == "__main__":
84     m, n, T, cacheStr = int(sys.argv[1]), int(sys.argv[2]), int(sys.argv[3]), sys.argv[4]
85     if cacheStr == "LRU":
86         cache = LRUCache(m)
87     elif cacheStr == "FIFO":
88         cache = FIFOCache(m)
89     else:
90         Exception()
91
92     model(n, cache, T)
```
